




Gaussian Distribution



1-d Gaussian

$$\mathcal{N}(\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

Bean machine: drop ball with pins



2-d Gaussian

$$f_{\mathbf{X}}(x_1, \dots, x_k) = \frac{\exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)}{\sqrt{(2\pi)^k |\boldsymbol{\Sigma}|}}$$


From wikipedia and <http://home.dei.polimi.it>

Gaussian Mixture Model

Assumptions

- Each data point comes from one of K classes.
- The cluster prior distribution w_j is *unknown*.
- Each class c_j follows a Gaussian

distribution:
$$P(x|c_j, \theta_j) = \frac{1}{\sqrt{2\pi\sigma_j^2}} e^{-\frac{(x-\mu_j)^2}{2\sigma_j^2}}$$



- The parameters for each class μ_j, σ_j are *unknown*(need to be learned).
- The probability of x_i is the sum over all classes, $P(x_i|\theta) = \sum_{j=1}^K P(x_i|c_j, \theta_j)P(c_j)$

Soft Clustering with Gaussian Mixture Model

- Every object i is assigned to one cluster j with a probability
 - $P(z_i = j) \in [0,1]$ and $\sum_j P(z_i = j) = 1$
 - Where z_i is a hidden variable of which cluster x_i belongs to.


Assume the parameters of the GMM have been learned

- The probability of x_i belonging to cluster c_j :

$$\begin{aligned} \square \quad P(z_i = c_j | x_i) &\propto P(x_i, z_i = c_j) \\ &= w_j P(x_i | z_i = c_j) \end{aligned}$$

Cluster prior probabilities

Probability density function of each cluster



The E-M(Expectation Maximization) Algorithm

- A framework to approach maximum likelihood or maximum a posteriori estimates of parameters in statistical models.
- Expectation Step:
 - Assigns objects to clusters according to the current soft clustering or parameters of probabilistic clusters
 - $w_{ij}^{t+1} = P(z_i = j | x_i, \theta_j^t) \propto w_j P(x_i | z_i = j, \theta_j^t)$ ←

Joint probability of x_i and its cluster c_j

- Maximization Step:
 - finds the new parameters of each cluster that maximize the expected likelihood
 - $\theta_{t+1} = \operatorname{argmax}_{\theta} \sum_i \sum_j w_{ij}^{t+1} \log L(x_i, z_j | \theta)$

Example: Applying E-M algorithm to 1-D GMM

- Iteratively do the following two steps

- E-Step: Evaluate the soft clustering probability according to $\mu_j^t, \sigma_j^t, w_j^t$

- $$w_{ij}^{t+1} = \frac{w_j^t P(x_i | \mu_j^t, \sigma_j^t)}{\sum_k w_k^t P(x_i | \mu_k^t, \sigma_k^t)}$$

- M-Step: Find the new parameters μ_j^t, σ_j^t that maximize log likelihood. In Gaussian distribution, this is equivalent to do parameter estimation when each data point has a weight.

- $$\mu_j^{t+1} = \frac{\sum_i w_{ij}^{t+1} x_i}{\sum_i w_{ij}^{t+1}}, (\sigma_j^2)^{t+1} = \frac{\sum_i w_{ij}^{t+1} (x_i - \mu_j^{t+1})^2}{\sum_i w_{ij}^{t+1}}$$




Weighted average means
and variance


- $$w_j^{t+1} = \frac{\sum_i w_{ij}^{t+1}}{n}$$

Gaussian Mixture Model


- Example of applying Gaussian Mixture Model



The data points belong to three classes. Each class follows a Gaussian distribution.



We hide the class information.



Class information inferred by GMM.

- We can use E-M algorithm to learn the parameters.

Example: Applying E-M algorithm to GMM

